



Matched-KL frontier — held-out lexical suppression per unit of distribution damage

Bands: 95% cluster-bootstrap CI (decisive pair only). n=200 held-out cases per model; oracle edits the evaluation tokens directly and is a labeled upper bound, not a label-free competitor.

Qwen3.5-2B (display model — replacement top-1 0.887)



R1-Distill-Qwen-7B (mid-fidelity readout — top-1 0.760)



SRP readout features mean-row direction PCA rank-1 PCA rank-4 oracle (uses eval labels) random features